

Locality Optimization for traversal-based Queries on Graph Databases

Fabian Klopfer

Databases and Information Systems Group
Department of Computer and Information Science
University of Konstanz, 30.04.2021


Motivation

Current state of performance-optimized

- relational databases: accesses are made as sequential as possible.
- graph databases: access is often random.

Example I

Show me all boats owned by sailors with an ID less than 50:

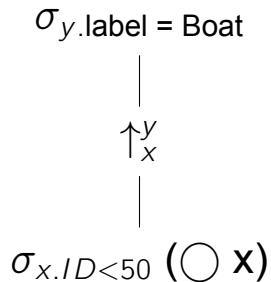


Reads are mostly sequential.

⇒ Prefetch & cache hit.

Example II

Nodes are Sailors and Boats, relationships “owns”



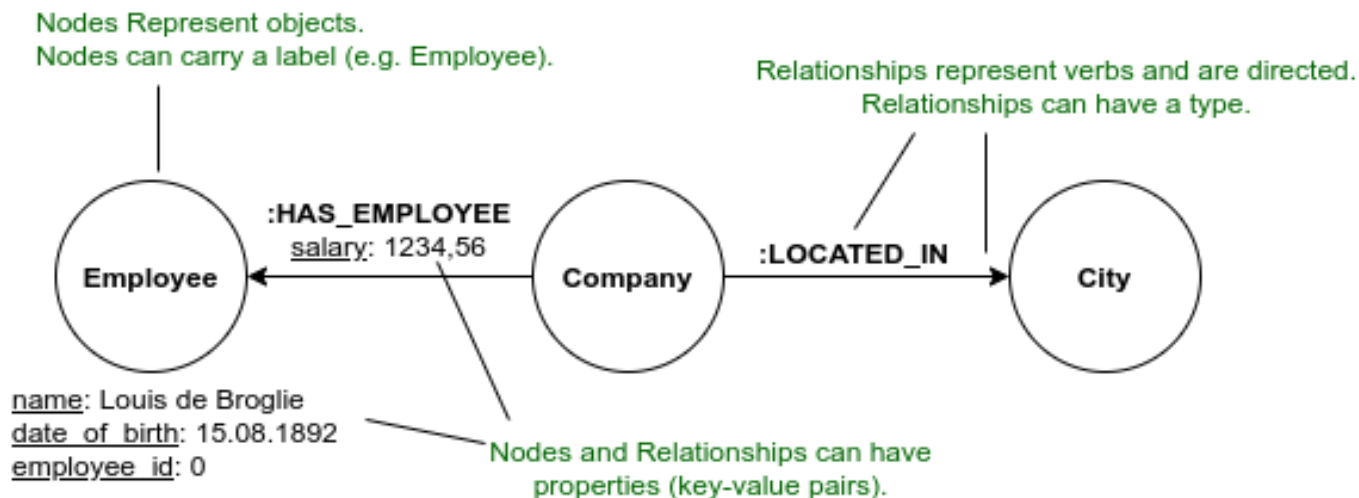
Scanning and filtering is sequential. Expand is not.

\Rightarrow Expand causes prefetch & cache misses.

Example III

- Especially `Expand` jumps a lot. Potentially back and forth.
- Traversals rely primarily on `expand`.

Property Graph Model



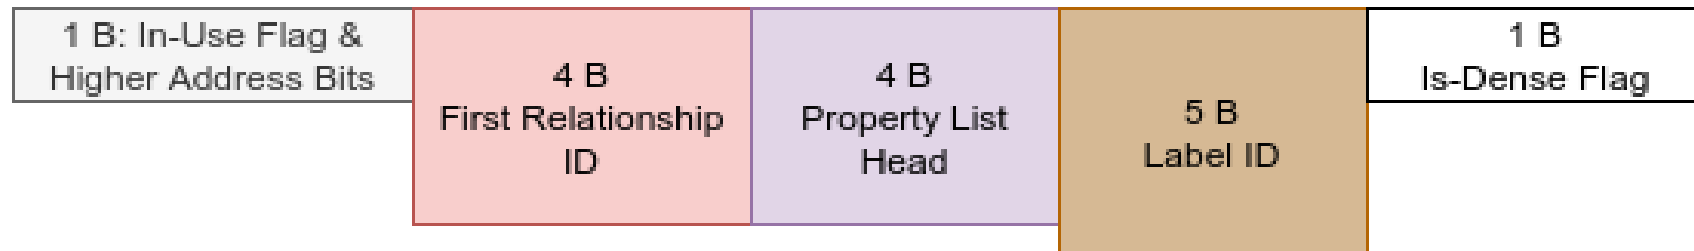
Data Structures I

Two essential record structures:


1. Node records
2. Relationship records

Inspired by Neo4J

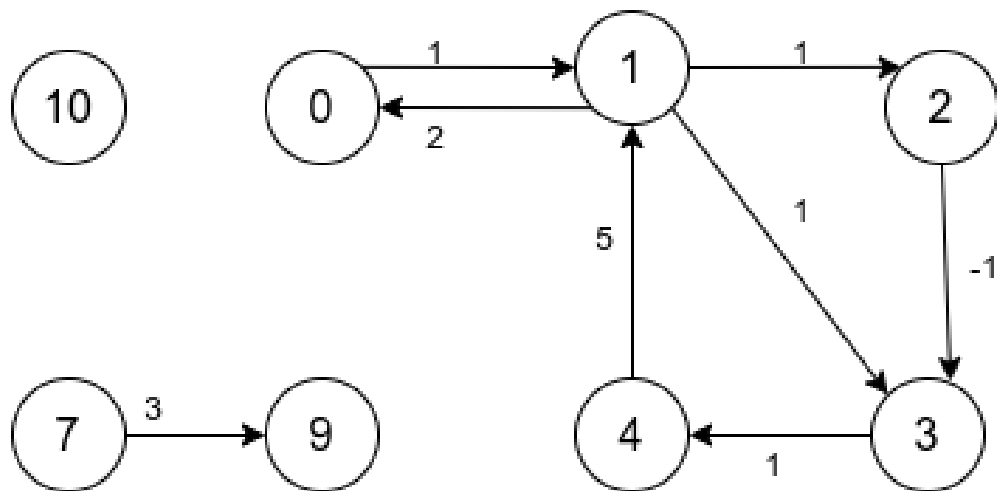
Data Structures II






Data Structures III



Graphs we focus on



Traversals



Problem Definition I

Given a graph G , logical block size b , page size p .

Desired is

1. A partition of G into blocks of vertex records V_i and E_i relationship records,
2. permutations π_v, π_e of the blocks of vertex and edge records V_i, E_i ,
3. a reordering of the incidence list pointers

such that spatial locality is as high as possible for traversal-based queries.

Problem Definition II


Temporal locality based on blocks.

$$P(X_{t+\Delta} = B | X_t = B)$$

Spatial locality in the same sense:

$$P(X_{t+\Delta} = B \pm \varepsilon | X_t = B)$$


Problem Definition III




node.db	0, 5, 7	1, 4, 9	2, 6, 8	3		
edge.db	a, f	b, g	c, h	d, i	e, j	k

node.db	7, 8, 9	0, 1, 3	2, 4, 5	6		
edge.db	f, h	g, k	i, j	a, b	c, d	e

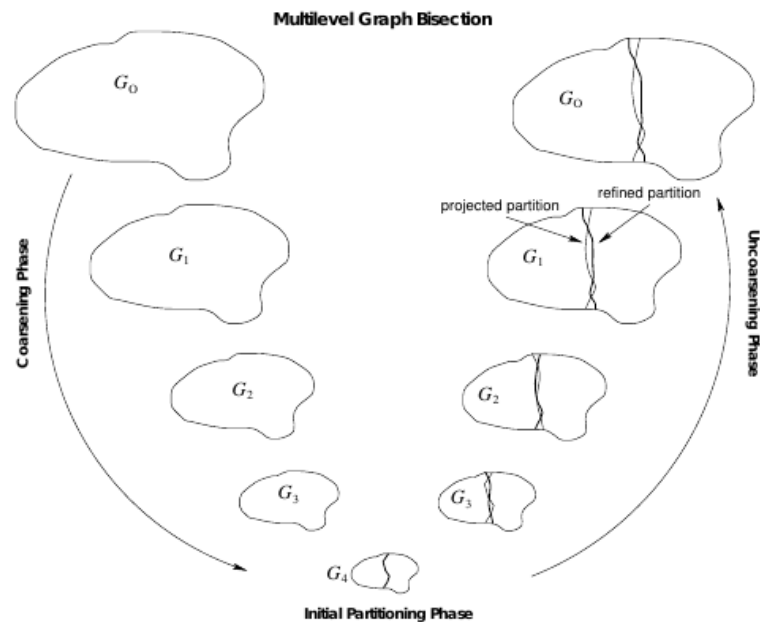
Problem Definition IV



3 Problem Definition



G-Store I

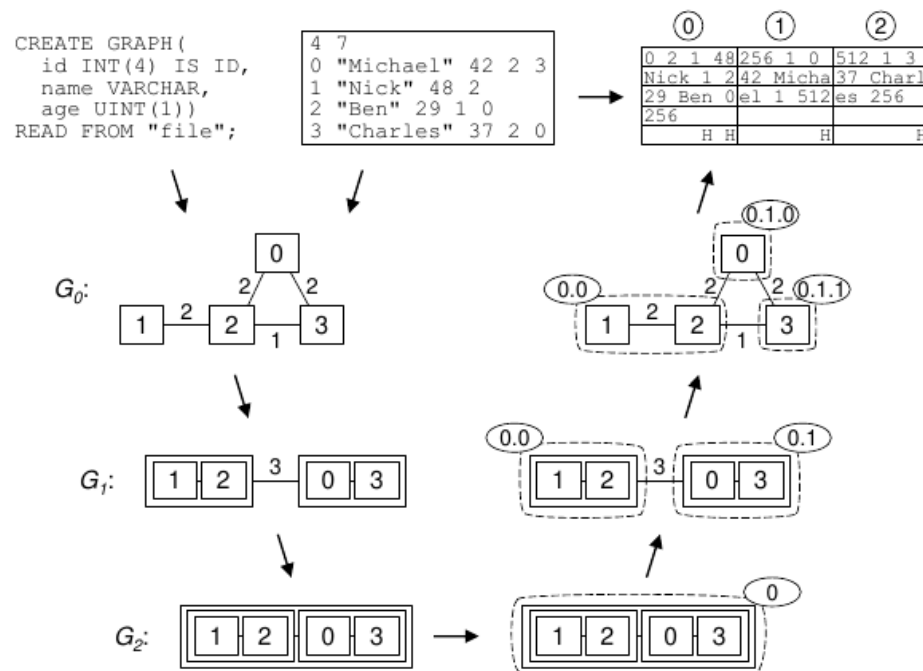


G-Store II

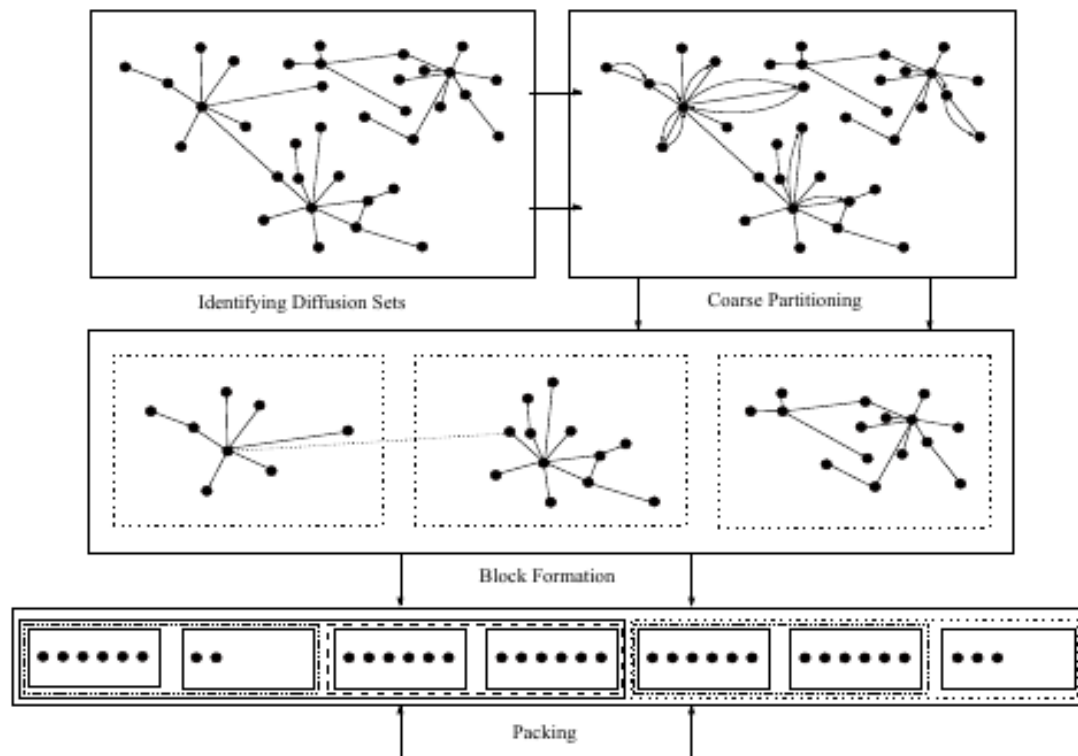
1. Coarsening: Heavy-Edge Matching
2. Turn-around
3. Uncoarsening
 - 3.1 Project
 - 3.2 Reorder
 - 3.3 Refine

$$\min \sum_{(u,v) \in E} |\phi(u) - \phi(v)|$$

G-Store III



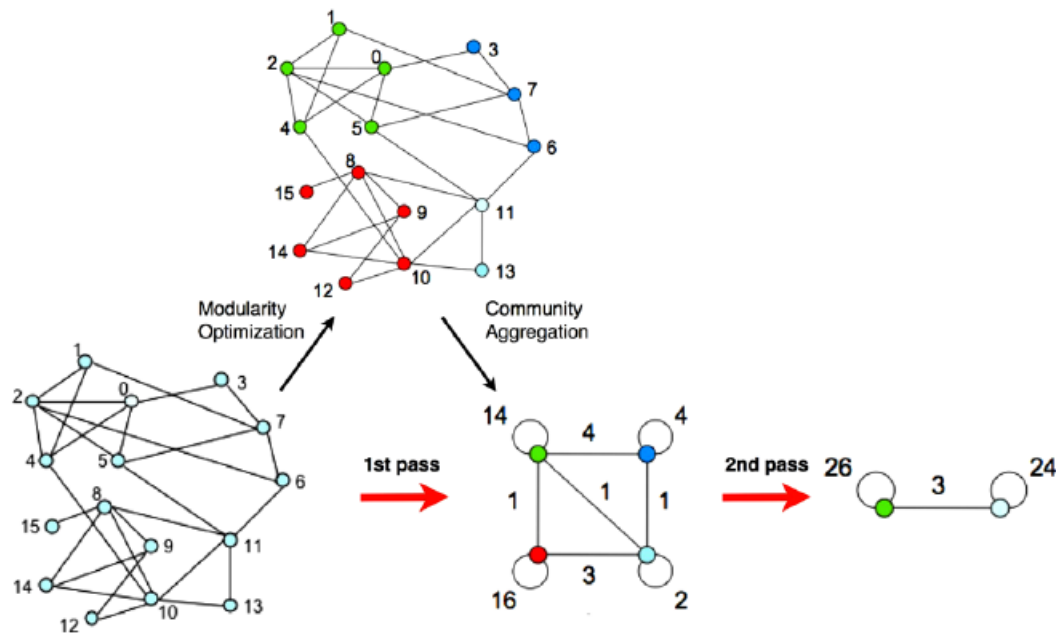
ICBL I



ICBL II

- I Feature extraction: Do t random walks of length l .
- C Coarse clustering: Adapted K-Means.
- B Block Formation: Agglomerative hierarchical clustering.
- L Layout Blocks: Sort blocks and subgraphs

Louvain Method I




Louvain Method II

1. Initialize all nodes in singleton community.
2. Merge community into a neighboring community where modularity gain is maximal, until modularity gain is below threshold.
3. Construct new graph from aggregated communities and go to 1.

$$\frac{1}{2m} \sum_{u,v \in V} \left(w_{(u,v)} - \frac{w_u w_v}{2m} \right) \cdot \delta(c_u, c_v)$$

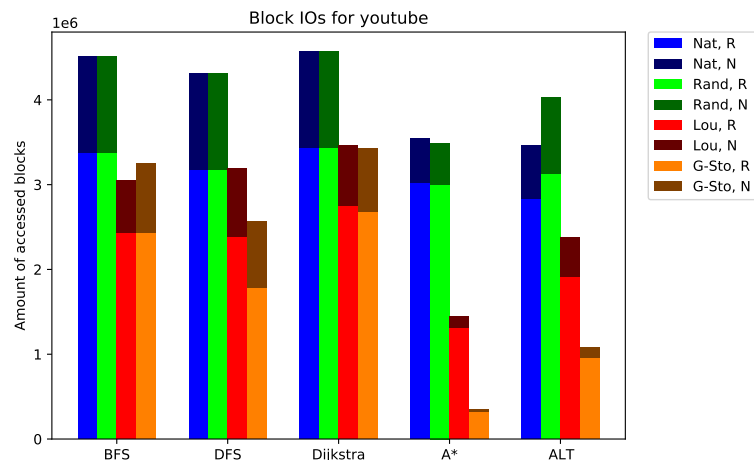
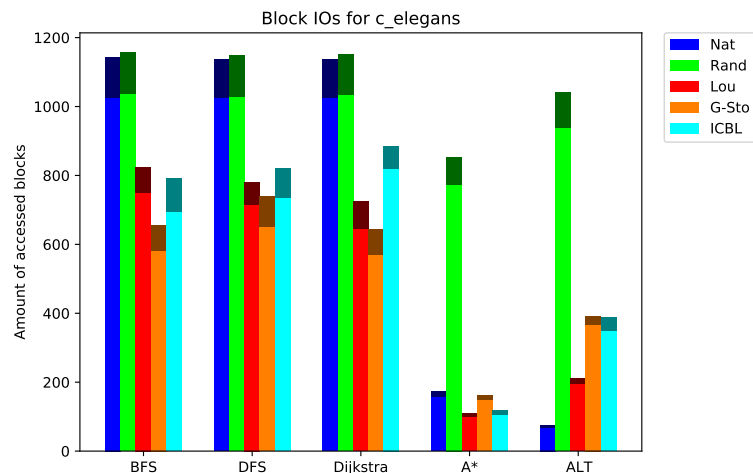
Incidence List Rearrangement




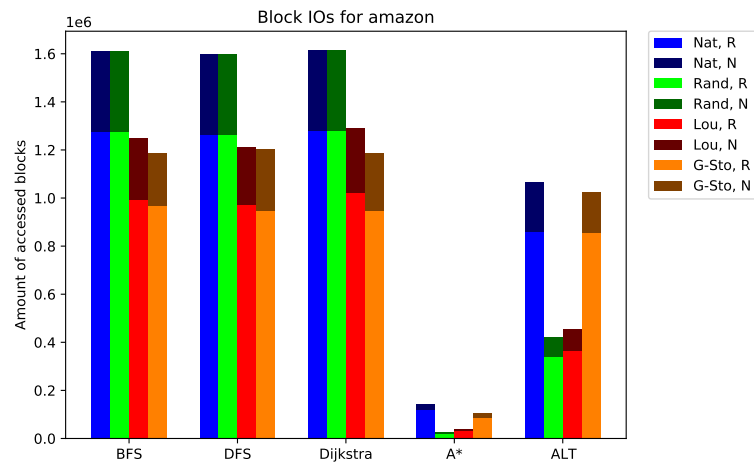
Setup

- Queries: BFS, DFS, Dijkstra, A^* , ALT.
- Datasets: [131, 1'134'890] nodes, [764, 2'987'624] edges, average degree [2.6, 25.5]
- Domains include biological neural net, E-Mails, Co-authors, Frequent item sets, Comments.

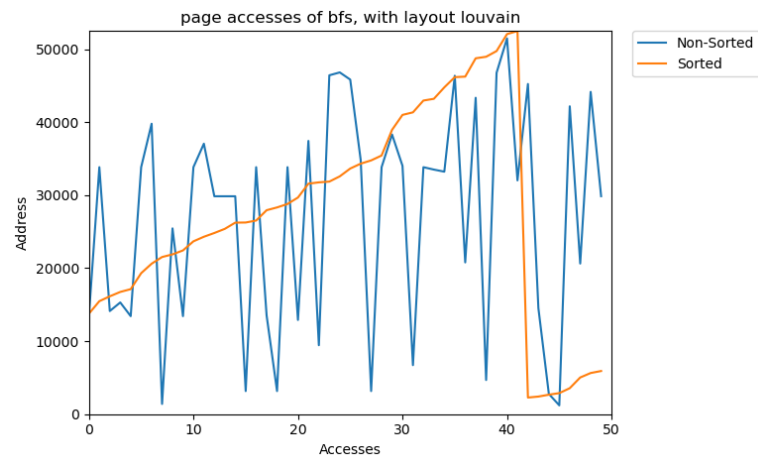
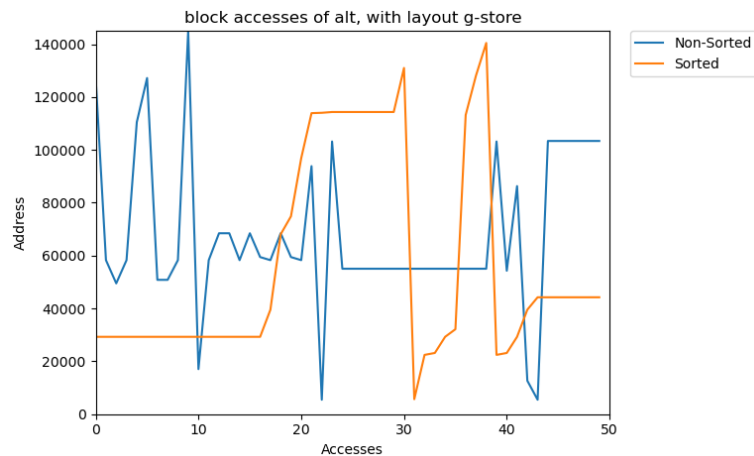
Results I



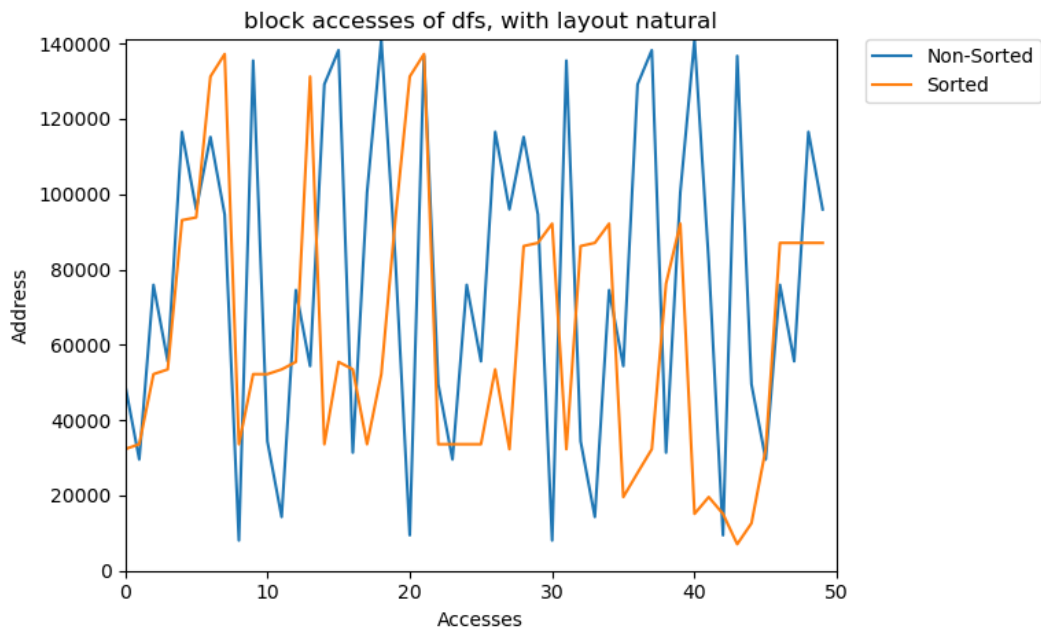
Results II



Results III



Results IV



Summary

- Static rearrangement methods decrease number of block accesses.
⇒ increase locality
- Sorting the incidence lists leads to more sequential access sequences.
- Ordering the blocks is crucial for spatial locality.

Future Work

- Leiden instead of Louvain
- RCM-based rearrangement
- Dynamic Rearrangement — Query-based