

# Locality Optimization for traversal-based Queries on Graph Databases

**Fabian Klopfer**

Databases and Information Systems Group  
Department of Computer and Information Science  
University of Konstanz, 30.04.2021

---

# Outline

- 1 Introduction
- 2 Background
- 3 Problem Definition
- 4 Locality-optimizing Record Layout
- 5 Evaluation
- 6 Conclusion


## Motivation I

### Current state of performance-optimized

- relational databases:
  - Storage order determined by join attribute. → enables localized access.
  - Explicit clustered indices (often B+-Trees). → represents order, speeds up range queries⇒ Accesses are made as sequential as possible.
- graph databases:
  - Storage order determined by insertion order.
  - Implicit, possibly unclustered index for relationships (doubly-linked list).
  - Lucene-based indexes on properties, labels, . . . , unclustered.⇒ Access is mostly random.

## Example I

Show me all boats owned by sailors with an ID less than 50:




Reads are mostly sequential.

$\Rightarrow$  Prefetch & cache hit.

## Example II

Nodes are Sailors and Boats, relationships “owns”




Scanning and filtering is sequential. Expand is not.

$\Rightarrow$  Expand causes prefetch & cache misses.



## Example III

- Especially `Expand` jumps a lot. Potentially back and forth.
- Traversals rely primarily on `expand`.

## Graphs

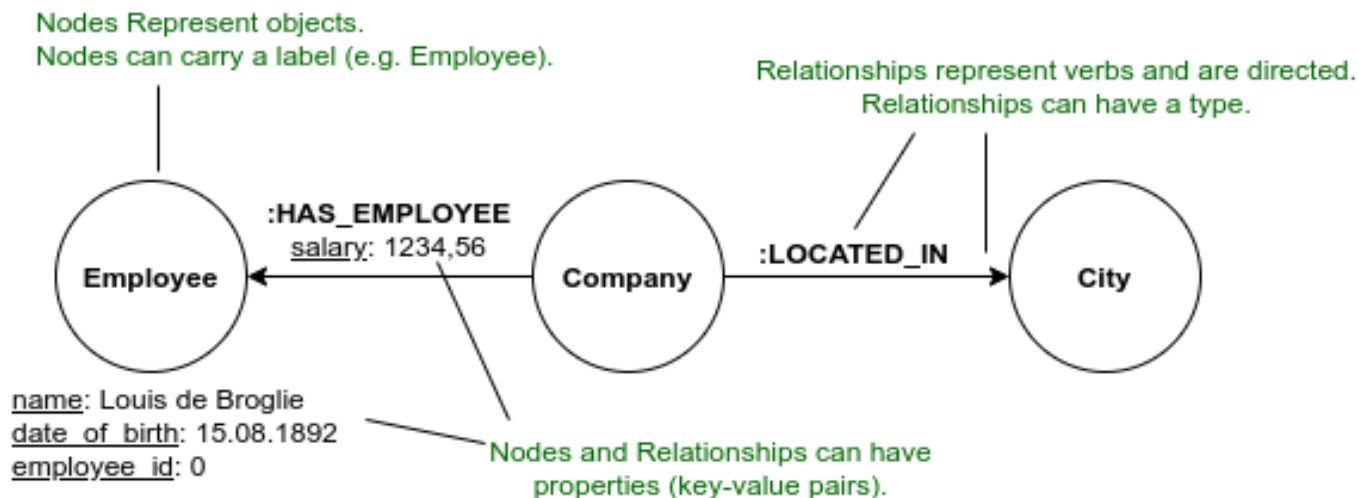


## Traversals





## Property Graph Model




## Data Structures I

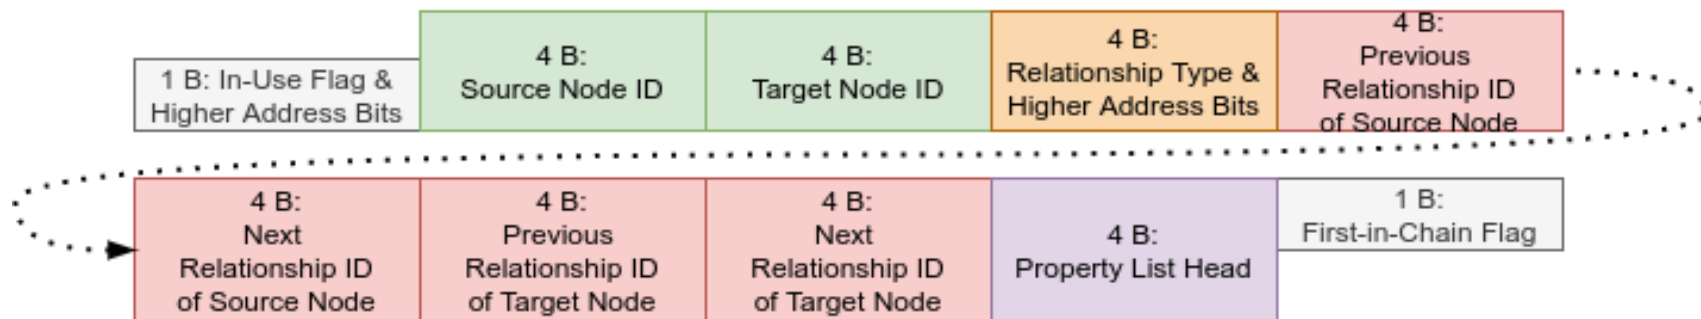
Two essential record structures:

1. Node records
2. Relationship records


## Data Structures II




## Data Structures III



## Data Structures IV



## 2 Background - 2.2 The



## Problem Definition I

Given a graph  $G$ , logical block size  $b$ , page size  $p$ .

Desired is

1. A partition of  $G$  into blocks of vertex records  $V_i$  and  $E_i$  relationship records,
2. permutations  $\pi_v, \pi_e$  of the blocks of vertex and edge records  $V_i, E_i$ ,
3. a reordering of the incidence list pointers

such that spatial locality is as high as possible for traversal-based queries.

## Problem Definition II

Temporal locality based on blocks.


$$P(X_{t+\Delta} = B | X_t = B)$$

Spatial locality in the same sense:

$$P(X_{t+\Delta} = B \pm \varepsilon | X_t = B)$$



# Problem Definition III




|         |         |         |         |      |      |   |
|---------|---------|---------|---------|------|------|---|
| node.db | 0, 5, 7 | 1, 4, 9 | 2, 6, 8 | 3    |      |   |
| edge.db | a, f    | b, g    | c, h    | d, i | e, j | k |

|         |         |         |         |      |      |   |
|---------|---------|---------|---------|------|------|---|
| node.db | 7, 8, 9 | 0, 1, 3 | 2, 4, 5 | 6    |      |   |
| edge.db | f, h    | g, k    | i, j    | a, b | c, d | e |

## Problem Definition IV

TODO insert incidence list chaos and ordered here


# G-Store I



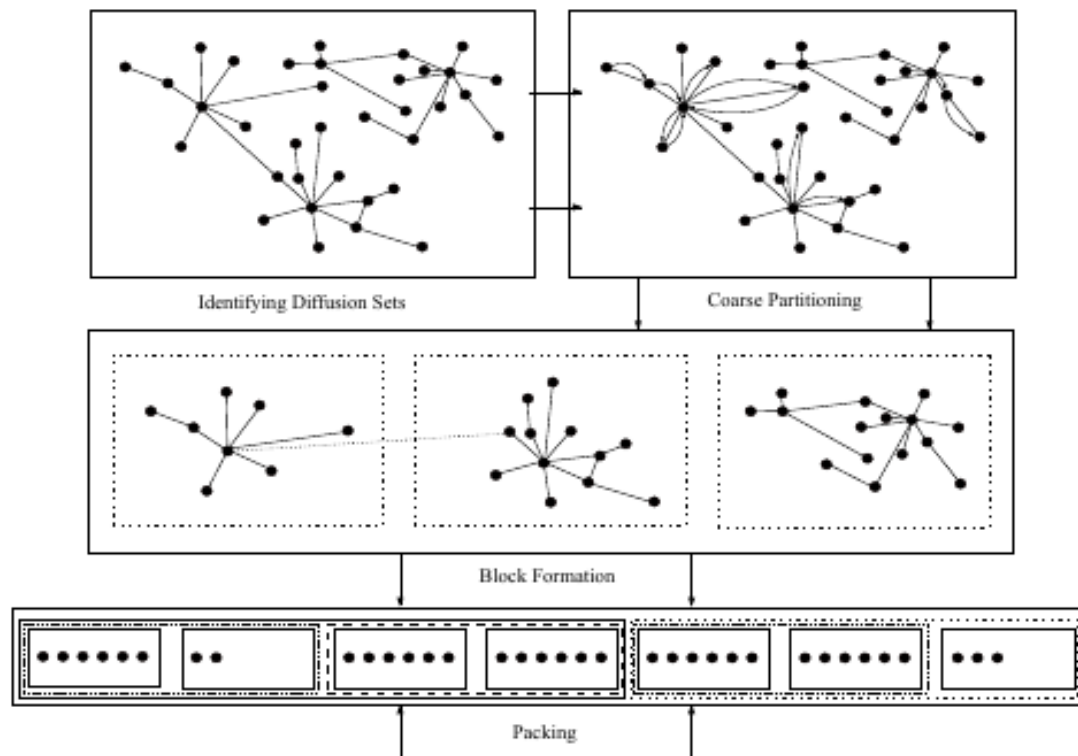
## G-Store II

1. Coarsening: Heavy-Edge Matching
2. Turn-around
3. Uncoarsening
  - 3.1 Project
  - 3.2 Reorder
  - 3.3 Refine

# G-Store III




### ICBL I



# ICBL II

- I Feature extraction: Do  $t$  random walks of length  $l$ .
- C Coarse clustering: Adapted K-Means.
- B Block Formation: Agglomerative hierarchical clustering.
- L Layout Blocks: Sort blocks and subgraphs

# Louvain Method I





## Louvain Method II

1. Initialize all nodes in singleton community.
2. Merge community into a neighboring community where modularity gain is maximal, until modularity gain is below threshold.
3. Construct new graph from aggregated communities and go to 1.

$$Q = \frac{1}{2m} \sum_{u,v \in V} \left( w_{(u,v)} - \frac{w_u w_v}{2m} \right) \cdot \delta(c_u, c_v)$$

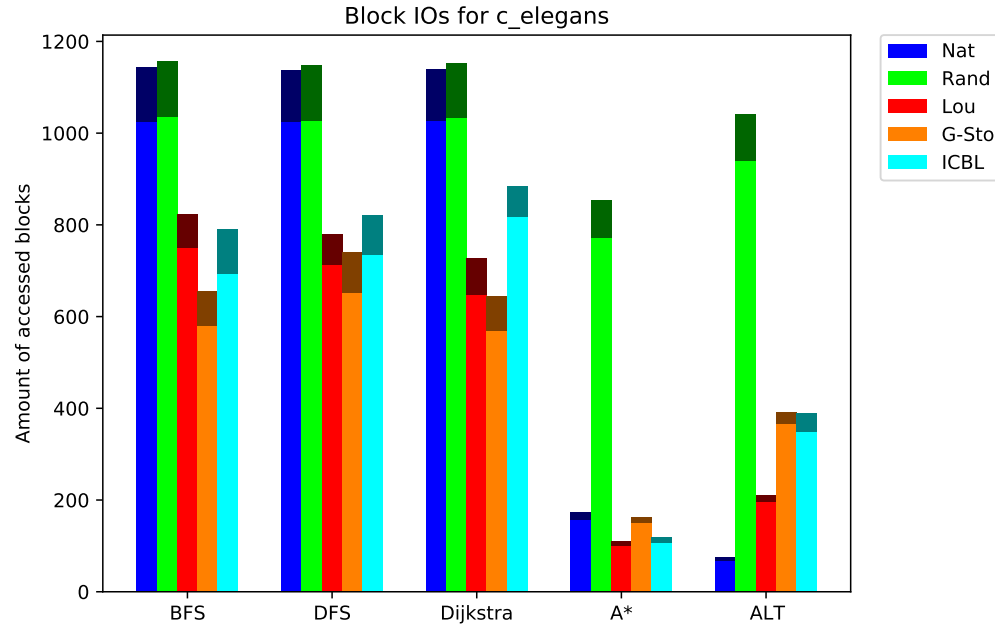
## Incidence List Rearrangement

1. Traverse incidence list and store IDs.
2. Sort IDs.
3. Assign first relationship pointer to lowest ID.
4. Assign next pointer of new first relationship to second ID.
5. Assign next pointer of relationship  $i$  to  $i + 1$ . ID and prev pointer to  $i - 1$ . ID.
6. Assign next pointer of last relation to first and prev pointer of first to last.

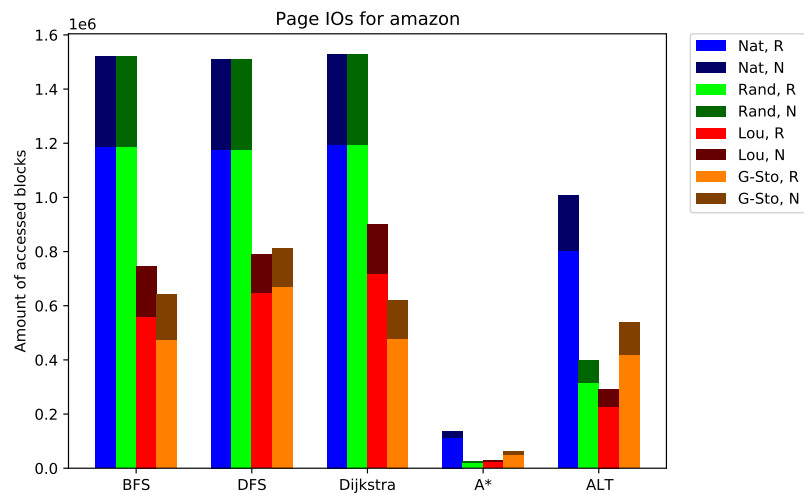

### Setup

- Queries: BFS, DFS, Dijkstra,  $A^*$ , ALT.
- Datasets: [131, 1'134'890] nodes, [764, 2'987'624] edges, average degree [2.6, 25.5]
- Domains include biological neural net, E-Mails, Co-authors, Frequent item sets, Comments.


# Results I




# Results II




## Results III



## Results IV

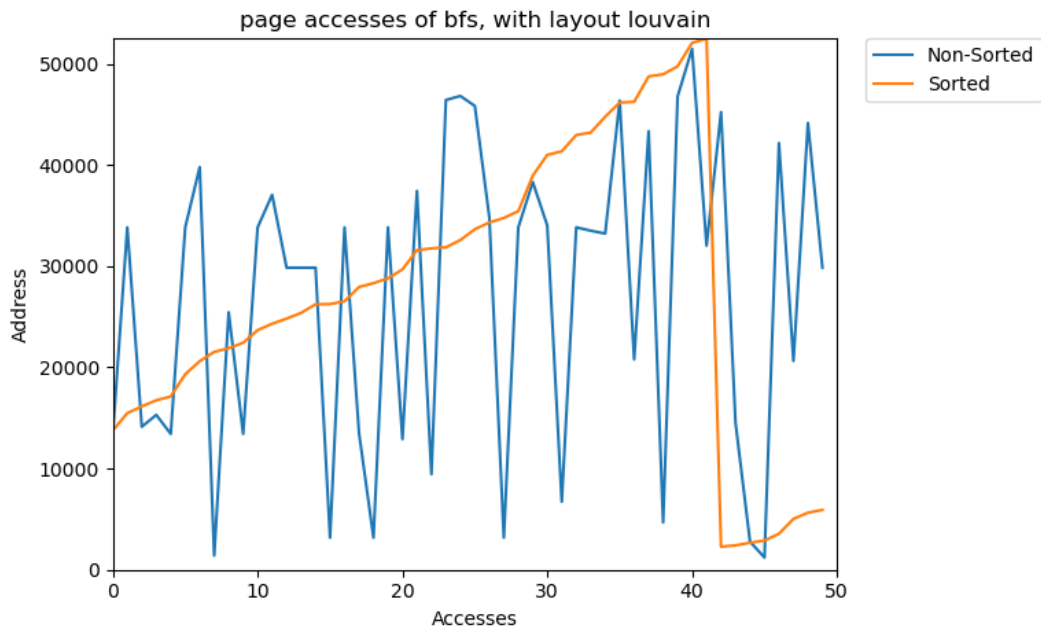


## Results V






# Results VI



## Results VII



# Summary I

- Static rearrangement methods decrease number of block accesses.  
⇒ increase locality
- Sorting the incidence lists leads to more sequential access sequences.
- Ordering the blocks is crucial for spatial locality.

# Future Work I

- Leiden instead of Louvain
- RCM-based rearrangement
- Dynamic Rearrangement — Query-based
- Disk-based implementation